

Word Embeddings

and its application in legal data feature extraction

Xu Xiao

Apr 15, 2019

Representation of linguistic data

QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

2

Hierarchy

- Document
- Paragraph
- Sentence
- Words



Vector Space

- Represent an item (e.g. a word) with a vector (list) of numbers.
- Computers can't recognize a word



Word vectors

QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

3

	emotional strength	positiveness	...
good	2	1	...
bad	2	0	...
great	4	1	...
terrible	4	0	...

More dimension means more information

- The properties of words can be endless
- Cannot generate automatically
- Different types of words require different properties

Representation of linguistic data

QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

4

“Words that occur in the same contexts tend to have similar meanings.”

(Harris, 1954)

“You shall know a word by the company it keeps.”

(Firth, 1957)

Counting over context

QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

5

	can	doesn't	hurt	...
eat	1	0	0	...
glass	0	0	0	...
it	0	1	0	...
me	0	0	1	...

Raw counting method uses numbers of occurrence as vector value

Given the sentence *I can eat glass and it doesn't hurt me*, suppose we take 2 adjacent words as context for each word, the vector space is shown to the left.

- Vector dimension equals to the length of vocabulary
- The vectors are sparse
- Frequent words may characterize vectors

tf-idf

QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

6

$$\text{tf}(t, d) = f_{t,d} / \sum_{t' \in d} f_{t',d}$$

$$\text{idf}(t, D) = \log \frac{N}{n_t} = -\log \frac{n_t}{N}$$

$$\text{tfidf}(t, d, D) = \text{tf}(t, d) \cdot \text{idf}(t, D)$$

Tf-idf adjust the weights with frequency in document

- Vector dimension equals to the length of vocabulary
- The vectors are sparse
- Frequent words may characterize vectors

We need to embedded the information into lower-dimension word vectors.

Word embeddings


QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

7



The low-dimension and dense word vectors are called word embeddings.

It involves a mathematical embedding from a space with one dimension per word to a continuous vector space with a much lower dimension.

Several famous models to produce embeddings:

- word2vec
- GloVe
- LDA

word2vec


QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

8



word2vec uses a Neural Network with a “bottleneck” to achieve lower dimension

The Neural Network takes word and content as input and output, and has a hidden layer with set number of nodes.

Mikolov, T., Sutskever, I., Chen, K., Corrado, G. S., & Dean, J. (2013). Distributed representations of words and phrases and their compositionality. In *Advances in neural information processing systems* (pp. 3111-3119).

A brief review of Neural Network


QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

9



A neural network consists of “neurons” which is based on neurons in human brain.

A neural network is an interconnected assembly of simple processing elements, units or nodes, whose functionality is loosely based on the animal neuron. The processing ability of the network is stored in the interunit connection strengths, or weights, obtained by a process of adaptation to, or learning from, a set of training patterns.

Example: single layer perceptron


QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

10



Suppose we have a simple network and we want to solve 'OR' problem

We set learning rate to 0.5, initial weights $w_1=w_2=w_b=0.5$.

Example: single layer perceptron


QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

11



x_1	x_2	w_1	w_2	w_b	y	t	$a(t-y)$
0	0	0.5	0.5	0.5	0.5	0	-0.25

Example: single layer perceptron


QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

12



x_1	x_2	w_1	w_2	w_b	y	t	$a(t-y)$
0	0	0.5	0.5	0.5	0.5	0	-0.25
1	0	0.5	0.5	0.25	0.75	1	0.125

Example: single layer perceptron


QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

13



x_1	x_2	w_1	w_2	w_b	y	t	$a(t-y)$
0	0	0.5	0.5	0.5	0.5	0	-0.25
1	0	0.5	0.5	0.25	0.75	1	0.125
0	1	0.625	0.5	0.375	0.875	1	0.0625

Example: single layer perceptron


QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

14



x_1	x_2	w_1	w_2	w_b	y	t	$a(t-y)$
0	0	0.5	0.5	0.5	0.5	0	-0.25
1	0	0.5	0.5	0.25	0.75	1	0.125
0	1	0.625	0.5	0.375	0.875	1	0.0625
1	1	0.625	0.5625	0.4375	1.625	1	0.3125

Structure of word2vec models


QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

15



Inputs: One-hot vectors (vectors with only one non-zero element, which is 1)

Example: *I can eat glass and it doesn't hurt me*

I: [1 0 0 0 0 0 0 0 0]

can: [0 1 0 0 0 0 0 0 0]

eat: [0 0 1 0 0 0 0 0 0]

...

Structure of word2vec models

QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

16

Source Text

The quick brown fox jumps over the lazy dog. →

The quick brown fox jumps over the lazy dog. →

The quick brown fox jumps over the lazy dog. →

The quick brown fox jumps over the lazy dog. →

Training Samples

(the, quick)
(the, brown)

(quick, the)
(quick, brown)
(quick, fox)

(brown, the)
(brown, quick)
(brown, fox)
(brown, jumps)

(fox, quick)
(fox, brown)
(fox, jumps)
(fox, over)

Two models: Skipgram vs CBOW

Skipgram: Using the middle word to predict its context

CBOW (continuous bag-of-words): Using the context to predict the middle word

Continuous bag-of words


QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

17



CBOW takes multiple one-hot vectors as input to predict middle word

Example: *I can eat **glass** and it doesn't hurt me*

Here the model takes 4 one-hot input vectors and we aim to have the output corresponds to *glass*. The values from weights matrix after training are the embeddings we're looking for.

Skipgram


QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

18



Skipgram takes the middle word as input to predict its context

Example: *I can eat **glass** and it doesn't hurt me*

Here the model takes an one-hot input vectors representing *glass* and we aim to have the output corresponds to its context. The values from weights matrix after training are the embeddings we're looking for.

Subsampling of Frequent Words

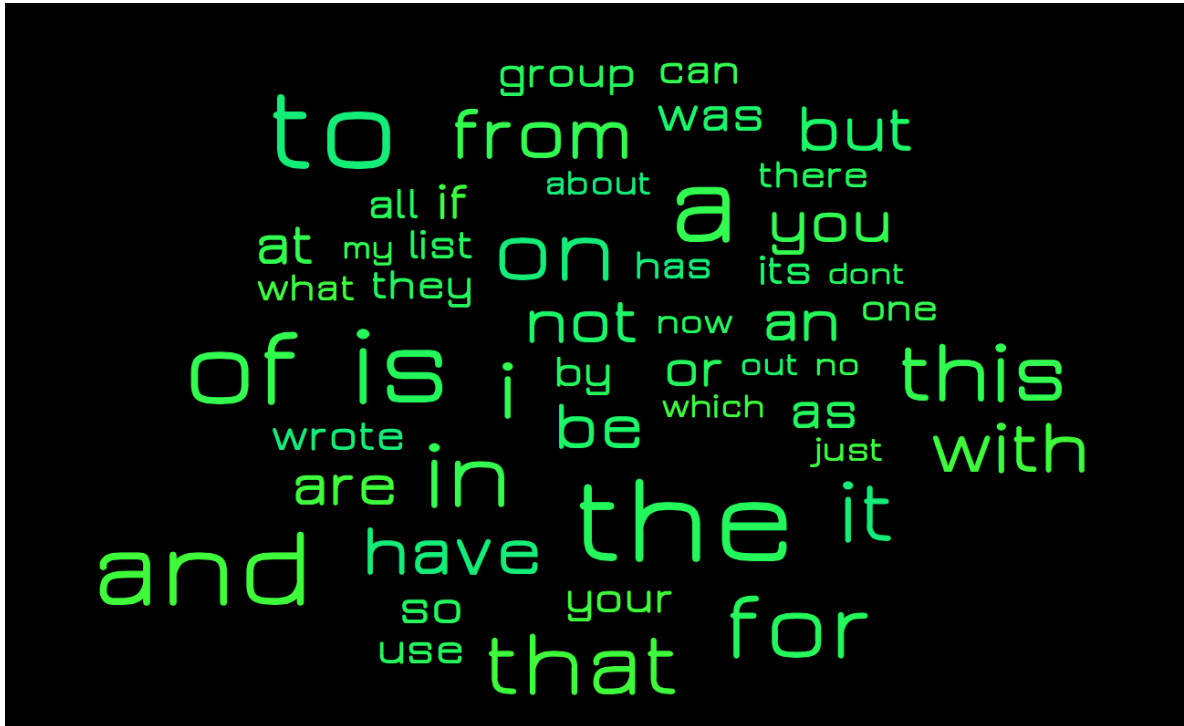
QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

19



To counter the imbalance between the rare and frequent words we need subsampling

Each word w_i in the training set is discarded with probability computed by the formula

$$P(w_i) = 1 - \sqrt{\frac{t}{f(w_i)}}$$

where $f(w_i)$ is the frequency of word w_i and t is a chosen threshold typically around 10^{-5} .

Using word2vec package

QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

20

Binary package

- <https://github.com/tmikolov/word2vec>
- Compile and download corpus
- Train your model



Using gensim

- Python package, install with pip
- Optimization over years
- Integration to your code
- <https://radimrehurek.com/gensim/models/word2vec.html>



Similarity


QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

21



One of the most direct application of word embeddings is to calculate similarity

The most common similarity measure of vectors is **Cosine similarity**.

$$\cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}},$$

TSNE graph

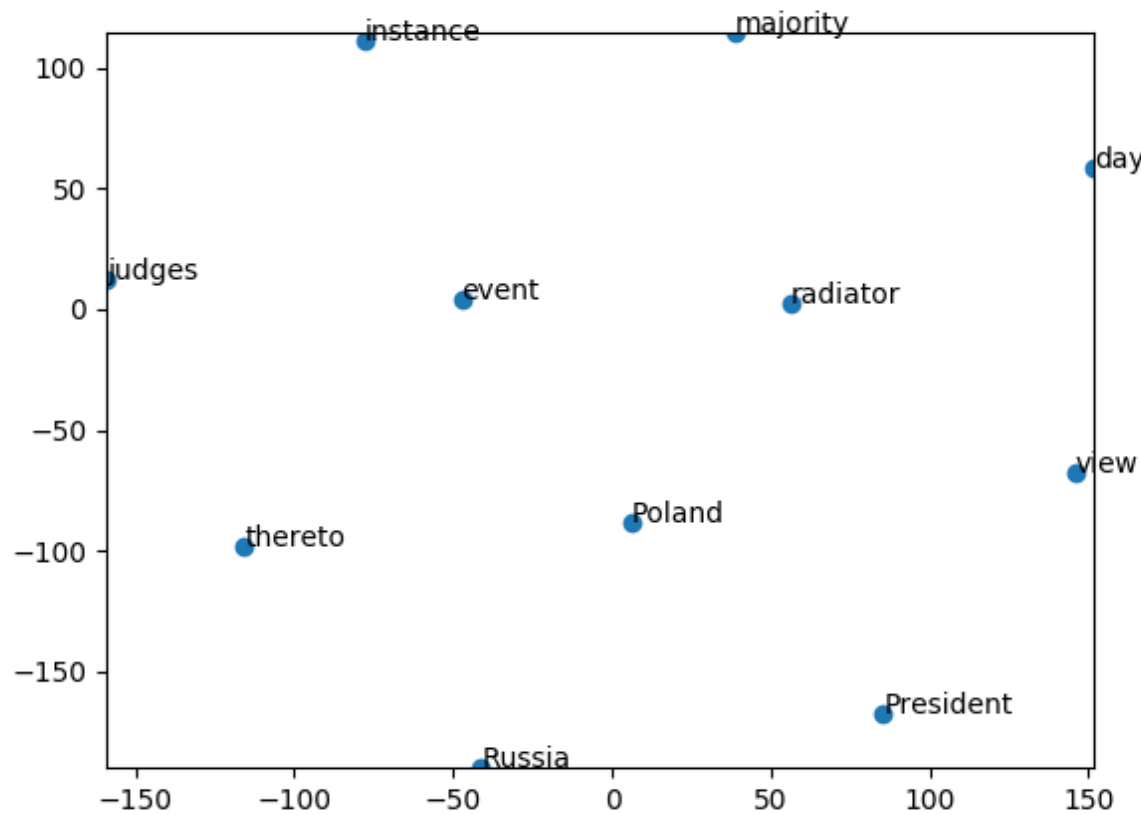
QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

22



You can use t-SNE graph to visualize high-dimension vectors

Visualizing Data using t-SNE, Maaten and Hinton, 2008

Analogies


QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

23



Word embeddings are capable of learning finding analogies

Suppose A, B, C and D are vectors representing words, If $B - A = D - C$, we could say the relations between AB and CD are similar.

Example:

$\text{vector}(\text{"king"}) - \text{vector}(\text{"man"}) = \text{vector}(\text{"queen"}) - \text{vector}(\text{"woman"})$

Feature extraction

QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION

24

Table 4: 10-fold cross-validation results: F1-score (macro) per article

Model	art3	art5	art6	art8
N-grams	0.78	0.68	0.61	0.50
Word embeddings	0.85	0.81	0.75	0.70

Word embeddings are usually combined with other classifiers

Using the vectors created for words in the word2vec dataset we determine the vectors for our training set and test set. However, as we do not want values per word, but per document, we average the vectors for each word in the entire document, and normalize them by using tf-idf weighting, which take into account the number of documents (i.e. cases) in which each word occurs.

Conclusions

QUESTIONS

INTRODUCTION

ALGORITHM

APPLICATION


25

Advantages

- A powerful word representation
- Incorporate well with other models

Disadvantages

- Need to custom parameters accordingly
- Need specified dataset
- Require large dataset and heavy computing



Questions

For lab session, please go to

<https://github.com/WillSkywalker/word-embedding>

Xu Xiao

Apr 15, 2019